

You don't choose a consumption option. You order them.

When a request comes back RESOURCE_EXHAUSTED, the reflex is to ask the same pool the same question again. A **capacity ladder** asks somewhere else, changing exactly one variable per step — consumption option, model or endpoint — until something answers. **No Provisioned Throughput order required.**

SIX STEPS, ONE VARIABLE EACH — THE DEFAULT, COST-EFFECTIVE LADDER

STEP	CONSUMPTION OPTION	MODEL	ENDPOINT	WHAT MOVED
1 flex	Flex PayGo	gemini-3.7-flash	global	—
2 standard	Standard PayGo	gemini-3.7-flash	global	tier
3 priority	Priority PayGo	gemini-3.7-flash	global	tier
4 alternative_model	Standard PayGo	gemini-3.6-flash	global	model
5 multi_region_us	Standard PayGo	gemini-3.6-flash	us	endpoint
6 multi_region_eu	Standard PayGo	gemini-3.6-flash	eu	endpoint

Each step changes exactly one variable — consumption option, model or endpoint.

- Steps 1-3 vary only the **consumption option**: Flex → Standard → Priority.
- Step 4 changes only the **model**.
- Steps 5-6 change only the **endpoint**.

CHOOSING A LADDER, OR MAKING YOUR OWN

PRESET	SHAPE	FOR
play_smart_default()	Flex → Standard → Priority → alternative → us → eu — the ladder above	Latency-tolerant, agentic
latency_first()	Priority on three models	Interactive, user-facing

Same engine, different order. A ladder is **data, not control flow** — reorder the list, or write your own in a dozen lines, and your cost and latency posture changes without touching client code.

THE 60-SECOND BET

Step 1 caps Flex at 60 seconds, one attempt, no retry. It is not a fallback — it is a **priced wager**: Flex is a published **50% discount**, so the step spends up to a minute of latency for a coin-flip at half price. The cap is the trick: **Flex's default timeout is ten minutes** — right for the offline work it is sold for, a stall as a synchronous first hop.

90

Tests that ship with it, offline and online — 81 need no cloud account and finish in under a second; 9 more drive the real ladder against **your own Google Cloud project**, including a drift check that tells you when the platform has moved. The 429 is real and free to summon: every attempt recorded, lane requested beside lane granted.

Read the article

The ladder step by step, the two headers, the 60-second Flex bet, and all 21 traversals with the attempt trails.

[Gemini Play Smart — a 429-survival capacity ladder →](#)

Run it against your project

Your entitlements, your quota, your regions. The 429 is real and free to summon — no simulation.

```
git clone https://github.com/eilonbar/gemini-play-smart
cd gemini-play-smart

python3 run.py
```

[github.com/eilonbar/gemini-play-smart →](#)

Disclaimer. This code is provided "as-is" as a demonstration only to illustrate a potential solution. The code does not constitute a Google product or service of any kind, and Google offers no support, warranties, or liability of any kind with its regard. Whoever chooses to use this code accepts all responsibility related to it, including for its implementation, use, and ongoing maintenance. For the avoidance of doubt, this code is not eligible for the Google Open Source Software Vulnerability Rewards Program.